

# Assignment 2 Report - Logistic Regression with $L_2$ Regularization

Helena Bales and Natalie Suderman

CS343 - Spring 2017

April 21, 2017

## **Contents**

# 1 Introduction

This report covers the second implementation assignment for Machine Learning and Data Mining. This assignment covers logistic regression with  $L_2$  Regularization using data from USPS's database of handwritten digits. This program will distinguish handwritten number 4's from 9's. Each digit is an image of size 16x16 pixels. The dataset that we will use was given on the class website and contains 700 training samples and 400 testing samples. We used the batch decent gradient algorithm to train to train the classifier. In order to complete this, we had to pick a reasonable learning rate and identify a reasonable stopping condition. These decisions will be covered in the Question 1 section below. Following completion of question 1, we reran the model from the beginning and plotted each decent gradient iteration. These plots, as well as an interpretation of the plots, can be found in the Question 2 section below. The next question involves  $L_2$  regularization. We will modify the code to find the gradient for the objective function  $L(w)$  given in the assignment and modify the batch gradient descent algorithm with this new gradient. The process and pseudocode for the  $L_2$  regularization will be covered in the Question 3 section below. Finally, the Question 4 section will cover the implementation of the change discussed in the Question 3 section. This implementation will then be evaluated based on their accuracies of the training and testing achieved using different  $\lambda$  values.

## 2 Question 1

### 2.1 Problem Statement

Implement the batch gradient descent algorithm to train a binary logistic regression classifier. The behavior of Gradient descent can be strongly influenced by the learning rate. Experiment with different learning rates, report your observation on the convergence behavior of the gradient descent algorithm. For your implementation, you will need to decide a stopping condition. You might use a fixed number of iterations, the change of the objective value (when it ceases to be significant) or the norm of the gradient (when it is smaller than a small threshold). Note, if you observe an overflow, then your learning rate is too big, so you need to try smaller learning rates.

### 2.2 Convergence Behavior

As the data was run multiple times the loss function was minimized. We ran about 1000 iterations, but in order to focus on the interesting parts, only the first 100 iterations were included in the interest of space.




Figure 1: ER diagram

## 2.3 Learning Rate

The learning rate we decided on was .00000001. Learning rates bigger than .0000001 always gave us an overflow error before the gradient converged.

## 2.4 Stopping Condition

Originally, we used the stopping condition  $\frac{d_{\text{New}} - d_{\text{Old}}}{d_{\text{Old}}} \geq .01$ , but the way that I calculated accuracy didn't provide enough significant digits for this to be successful. Instead we chose to stop using a set number of iterations. In our case, about 1000 iterations was sufficient to see the convergence.

# 3 Question 2

## 3.1 Problem Statement

Once you identify a suitable learning rate, rerun the training of the model from the beginning. For each gradient descent iteration, plot the training accuracy and the testing accuracy of your model as a function of the number of gradient descent iterations. What trend do you observe?

## 3.2 Method

## 3.3 Plots




Figure 2: ER diagram

## 3.4 Interpretation of Data

The test data immediately begins at a higher accuracy, because it has the advantage of starting with the weight vector calculated from the train data. You can see that after several iterations the model improves dramatically, and rests at about a 95% accuracy for both sets of data.

# 4 Question 3

## 4.1 Problem Statement

As discussed in class, Logistic regression is typically used with regularization. We will explore  $L_2$  regularization for this question. In particular, we will the following objective with an additional regularization term that is equal to the squared Euclidean norm of the weight vector.

Where the loss function  $l$  is the same as introduced in class (slide 7 of the logistic regression notes). Find the gradient for this objective function and modify the

batch gradient descent algorithm with this new gradient. Provide the pseudo code for your modified algorithm.

## **4.2 Method**

## **4.3 Pseudocode**

# **5 Question 4**

## **5.1 Problem Statement**

## **5.2 Implementation of Pseudocode from Question 3**

## **5.3 Evaluation of Implementation**

# **6 Conclusion**